

Loading Libraries

```
In [2]: import tensorflow as tf
        from tensorflow.keras import datasets, layers, models
        import matplotlib.pyplot as plt

        import warnings
        warnings.filterwarnings('ignore')
```

WARNING:tensorflow:From c:\Users\rodri\AppData\Local\Programs\Python\Python39\lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

Loading Data

```
In [3]: (train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data(

        # Normalize pixel values to be between 0 and 1 by dividing by 255
        train_images, test_images = train_images / 255.0, test_images / 255.0
```

Training the Models

Model 1: Artificial Neural Networks (ANNs)

In this model exploration, I aim to optimize the performance of an image classification model using Artificial Neural Networks (ANNs) by adjusting various configurations and hyperparameters. The initial focus will be on exploring and comparing different optimizer algorithms, which play a critical role in determining how the network learns from the data by adjusting its weights during the training process.

Optimizers

Optimizers in ANNs are pivotal in refining the model's parameters concerning a defined loss function, essentially guiding the learning process by updating the weights iteratively. For this investigation, I'll delve into three distinct optimizer algorithms:

- Adam (Adaptive Moment Estimation): This algorithm combines the advantages of both Adaptive Gradient Algorithm (AdaGrad) and Root Mean Square Propagation (RMSProp) approaches. It maintains adaptive learning rates for each parameter, adjusting them as the training progresses based on the moments of gradients.
- SGD (Stochastic Gradient Descent): This classic optimization algorithm updates the model's parameters proportional to the gradient of the error with respect to a single training example, often utilizing a learning rate to determine the step size in the parameter space.

- RMSProp (Root Mean Square Propagation): RMSProp adjusts the learning rates for different model parameters by dividing the learning rate for a weight by the average of the magnitudes of recent gradients for that weight.

```
In [4]: # Create ANN model with the Adam optimizer
ann = models.Sequential([
    layers.Flatten(input_shape = (32,32,3)),
    layers.Dense(128, activation = 'relu'),
    layers.Dense(64, activation = 'relu'),
    layers.Dense(10, activation = 'softmax')
])

ann.compile(optimizer = 'adam',
            loss = 'sparse_categorical_crossentropy',
            metrics = ['accuracy'])

# Train the model
history = ann.fit(train_images, train_labels, epochs=10,
                  validation_data=(test_images, test_labels))

# Evaluate the model
test_loss, test_acc = ann.evaluate(test_images, test_labels)
print(f"Test accuracy: {test_acc}")

# Plot training history
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
plt.show()
```

WARNING:tensorflow:From c:\Users\rodri\AppData\Local\Programs\Python\Python39\lib\site-packages\keras\src\backend.py:873: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From c:\Users\rodri\AppData\Local\Programs\Python\Python39\lib\site-packages\keras\src\optimizers__init__.py:309: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

Epoch 1/10

WARNING:tensorflow:From c:\Users\rodri\AppData\Local\Programs\Python\Python39\lib\site-packages\keras\src\utils\tf_utils.py:492: The name tf.ragged.RaggedTensorValue is deprecated. Please use tf.compat.v1.ragged.RaggedTensorValue instead.

WARNING:tensorflow:From c:\Users\rodri\AppData\Local\Programs\Python\Python39\lib\site-packages\keras\src\engine\base_layer_utils.py:384: The name tf.executing_eagerly_outside_functions is deprecated. Please use tf.compat.v1.executing_eagerly_outside_functions instead.

1563/1563 [=====] - 6s 3ms/step - loss: 1.8998 - accuracy: 0.3103 - val_loss: 1.7608 - val_accuracy: 0.3748

Epoch 2/10

1563/1563 [=====] - 5s 3ms/step - loss: 1.7240 - accuracy: 0.3814 - val_loss: 1.6405 - val_accuracy: 0.4197

Epoch 3/10

1563/1563 [=====] - 5s 3ms/step - loss: 1.6530 - accuracy: 0.4066 - val_loss: 1.6171 - val_accuracy: 0.4271

Epoch 4/10

1563/1563 [=====] - 5s 3ms/step - loss: 1.6082 - accuracy: 0.4256 - val_loss: 1.6199 - val_accuracy: 0.4244

Epoch 5/10

1563/1563 [=====] - 5s 3ms/step - loss: 1.5701 - accuracy: 0.4389 - val_loss: 1.5792 - val_accuracy: 0.4366

Epoch 6/10

1563/1563 [=====] - 5s 3ms/step - loss: 1.5429 - accuracy: 0.4506 - val_loss: 1.5529 - val_accuracy: 0.4449

Epoch 7/10

1563/1563 [=====] - 5s 3ms/step - loss: 1.5229 - accuracy: 0.4571 - val_loss: 1.5820 - val_accuracy: 0.4335

Epoch 8/10

1563/1563 [=====] - 5s 3ms/step - loss: 1.5026 - accuracy: 0.4642 - val_loss: 1.4944 - val_accuracy: 0.4674

Epoch 9/10


1563/1563 [=====] - 5s 3ms/step - loss: 1.4877 - accuracy: 0.4668 - val_loss: 1.4859 - val_accuracy: 0.4723

Epoch 10/10

1563/1563 [=====] - 5s 3ms/step - loss: 1.4752 - accuracy: 0.4722 - val_loss: 1.5135 - val_accuracy: 0.4651

313/313 [=====] - 0s 1ms/step - loss: 1.5135 - accuracy: 0.4651

Test accuracy: 0.4650999903678894



```
In [4]: # Create ANN model with the SGD optimizer
ann = models.Sequential([
    layers.Flatten(input_shape = (32,32,3)),
    layers.Dense(128, activation = 'relu'),
    layers.Dense(64, activation = 'relu'),
    layers.Dense(10, activation = 'softmax')
])


ann.compile(optimizer = 'sgd',
            loss = 'sparse_categorical_crossentropy',
            metrics = ['accuracy'])

# Train the model
history = ann.fit(train_images, train_labels, epochs=10,
                  validation_data=(test_images, test_labels))

# Evaluate the model
test_loss, test_acc = ann.evaluate(test_images, test_labels)
print(f"Test accuracy: {test_acc}")

# Plot training history
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
plt.show()
```

Epoch 1/10
 1563/1563 [=====] - 4s 2ms/step - loss: 1.8969 - accuracy: 0.3192 - val_loss: 1.7575 - val_accuracy: 0.3774
 Epoch 2/10
 1563/1563 [=====] - 3s 2ms/step - loss: 1.7102 - accuracy: 0.3954 - val_loss: 1.6654 - val_accuracy: 0.4160
 Epoch 3/10
 1563/1563 [=====] - 3s 2ms/step - loss: 1.6349 - accuracy: 0.4206 - val_loss: 1.6157 - val_accuracy: 0.4308
 Epoch 4/10
 1563/1563 [=====] - 3s 2ms/step - loss: 1.5811 - accuracy: 0.4393 - val_loss: 1.6265 - val_accuracy: 0.4189
 Epoch 5/10
 1563/1563 [=====] - 3s 2ms/step - loss: 1.5424 - accuracy: 0.4526 - val_loss: 1.5442 - val_accuracy: 0.4542
 Epoch 6/10
 1563/1563 [=====] - 3s 2ms/step - loss: 1.5121 - accuracy: 0.4638 - val_loss: 1.5240 - val_accuracy: 0.4554
 Epoch 7/10
 1563/1563 [=====] - 3s 2ms/step - loss: 1.4834 - accuracy: 0.4730 - val_loss: 1.5039 - val_accuracy: 0.4653
 Epoch 8/10
 1563/1563 [=====] - 3s 2ms/step - loss: 1.4561 - accuracy: 0.4849 - val_loss: 1.4865 - val_accuracy: 0.4673
 Epoch 9/10
 1563/1563 [=====] - 3s 2ms/step - loss: 1.4341 - accuracy: 0.4919 - val_loss: 1.4912 - val_accuracy: 0.4725
 Epoch 10/10
 1563/1563 [=====] - 3s 2ms/step - loss: 1.4139 - accuracy: 0.5008 - val_loss: 1.5458 - val_accuracy: 0.4609
 313/313 [=====] - 0s 949us/step - loss: 1.5458 - accuracy: 0.4609
 Test accuracy: 0.4609000086784363



```
In [5]: # Create ANN model with the RMSProp optimizer
ann = models.Sequential([
    layers.Flatten(input_shape = (32,32,3)),
    layers.Dense(128, activation = 'relu'),
    layers.Dense(64, activation = 'relu'),
    layers.Dense(10, activation = 'softmax')
])

ann.compile(optimizer = 'rmsprop',
            loss = 'sparse_categorical_crossentropy',
```

```

        metrics = ['accuracy'])

# Train the model
history = ann.fit(train_images, train_labels, epochs=10,
                  validation_data=(test_images, test_labels))

# Evaluate the model
test_loss, test_acc = ann.evaluate(test_images, test_labels)
print(f"Test accuracy: {test_acc}")


# Plot training history
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
plt.show()

```

```

Epoch 1/10
1563/1563 [=====] - 5s 3ms/step - loss: 1.9302 - accuracy:
0.2988 - val_loss: 1.7885 - val_accuracy: 0.3515
Epoch 2/10
1563/1563 [=====] - 4s 3ms/step - loss: 1.7426 - accuracy:
0.3777 - val_loss: 1.7004 - val_accuracy: 0.3867
Epoch 3/10
1563/1563 [=====] - 4s 3ms/step - loss: 1.6792 - accuracy:
0.4010 - val_loss: 1.6222 - val_accuracy: 0.4222
Epoch 4/10
1563/1563 [=====] - 4s 3ms/step - loss: 1.6374 - accuracy:
0.4152 - val_loss: 1.6799 - val_accuracy: 0.3989
Epoch 5/10
1563/1563 [=====] - 4s 3ms/step - loss: 1.6075 - accuracy:
0.4301 - val_loss: 1.6709 - val_accuracy: 0.4101
Epoch 6/10
1563/1563 [=====] - 4s 3ms/step - loss: 1.5862 - accuracy:
0.4381 - val_loss: 1.6372 - val_accuracy: 0.4208
Epoch 7/10
1563/1563 [=====] - 4s 3ms/step - loss: 1.5672 - accuracy:
0.4411 - val_loss: 1.6096 - val_accuracy: 0.4271
Epoch 8/10
1563/1563 [=====] - 4s 3ms/step - loss: 1.5550 - accuracy:
0.4475 - val_loss: 1.5519 - val_accuracy: 0.4541
Epoch 9/10
1563/1563 [=====] - 4s 3ms/step - loss: 1.5399 - accuracy:
0.4523 - val_loss: 1.6618 - val_accuracy: 0.4263
Epoch 10/10
1563/1563 [=====] - 4s 3ms/step - loss: 1.5319 - accuracy:
0.4570 - val_loss: 1.6914 - val_accuracy: 0.4221
313/313 [=====] - 0s 995us/step - loss: 1.6914 - accuracy:
0.4221
Test accuracy: 0.4221000075340271

```



Results

A 46.50% accuracy score with the Adam optimizer showcases its performance as the most effective among the three optimizers—Adam, SGD, and RMSProp—utilized in the image classification task. This accuracy metric reflects the model's ability to correctly classify nearly half of the images within the dataset. Adam's success in achieving the highest accuracy suggests its suitability for this specific dataset and neural network configuration. Its adaptive learning rate and combination of momentum techniques seem to have facilitated better convergence and parameter optimization compared to the other optimizers.

Layers

Now, the exploration shifts towards varying the number of layers within the neural network architecture. The number of layers profoundly influences the network's capacity to learn intricate patterns and features within the dataset.

This phase involves experimenting with different layer configurations, such as:

- **Shallow Networks:** Consisting of fewer layers, typically with a small number of hidden layers.
- **Deep Networks:** Incorporating a larger number of layers, allowing for a more intricate hierarchy of features and representations.

By systematically adjusting the number of layers while maintaining other hyperparameters constant or within defined ranges, the goal is to observe how the model's accuracy responds to these structural modifications. Each configuration will undergo training and evaluation, measuring accuracy metrics and potentially other performance indicators.

```
In [6]: # Create a Shallow Networks model
ann = models.Sequential([
    layers.Flatten(input_shape = (32,32,3)),
    layers.Dense(128, activation = 'relu'),
    layers.Dense(10, activation = 'softmax')
])
```

```

ann.compile(optimizer = 'adam',
            loss = 'sparse_categorical_crossentropy',
            metrics = ['accuracy'])

# Train the model
history = ann.fit(train_images, train_labels, epochs=10,
                  validation_data=(test_images, test_labels))

# Evaluate the model
test_loss, test_acc = ann.evaluate(test_images, test_labels)
print(f"Test accuracy: {test_acc}")


# Plot training history
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
plt.show()

```

```

Epoch 1/10
1563/1563 [=====] - 5s 3ms/step - loss: 1.9382 - accuracy:
0.2934 - val_loss: 1.8268 - val_accuracy: 0.3300
Epoch 2/10
1563/1563 [=====] - 5s 3ms/step - loss: 1.8108 - accuracy:
0.3431 - val_loss: 1.7695 - val_accuracy: 0.3636
Epoch 3/10
1563/1563 [=====] - 4s 3ms/step - loss: 1.7589 - accuracy:
0.3638 - val_loss: 1.7479 - val_accuracy: 0.3636
Epoch 4/10
1563/1563 [=====] - 4s 3ms/step - loss: 1.7315 - accuracy:
0.3747 - val_loss: 1.7122 - val_accuracy: 0.3850
Epoch 5/10
1563/1563 [=====] - 4s 3ms/step - loss: 1.7154 - accuracy:
0.3814 - val_loss: 1.7029 - val_accuracy: 0.3913
Epoch 6/10
1563/1563 [=====] - 4s 3ms/step - loss: 1.7053 - accuracy:
0.3843 - val_loss: 1.7167 - val_accuracy: 0.3800
Epoch 7/10
1563/1563 [=====] - 4s 3ms/step - loss: 1.6941 - accuracy:
0.3899 - val_loss: 1.7449 - val_accuracy: 0.3708
Epoch 8/10
1563/1563 [=====] - 4s 3ms/step - loss: 1.6869 - accuracy:
0.3915 - val_loss: 1.6915 - val_accuracy: 0.3855
Epoch 9/10
1563/1563 [=====] - 4s 3ms/step - loss: 1.6803 - accuracy:
0.3957 - val_loss: 1.6943 - val_accuracy: 0.3809
Epoch 10/10
1563/1563 [=====] - 4s 3ms/step - loss: 1.6765 - accuracy:
0.3949 - val_loss: 1.6778 - val_accuracy: 0.3961
313/313 [=====] - 0s 917us/step - loss: 1.6778 - accuracy:
0.3961
Test accuracy: 0.3961000144481659

```

```
In [7]: # Create a Deep Networks model
ann = models.Sequential([
    layers.Flatten(input_shape = (32,32,3)),
    layers.Dense(256, activation = 'relu'),
    layers.Dense(128, activation = 'relu'),
    layers.Dense(64, activation = 'relu'),
    layers.Dense(32, activation = 'relu'),
    layers.Dense(10, activation = 'softmax')
])


ann.compile(optimizer = 'adam',
            loss = 'sparse_categorical_crossentropy',
            metrics = ['accuracy'])

# Train the model
history = ann.fit(train_images, train_labels, epochs=10,
                  validation_data=(test_images, test_labels))

# Evaluate the model
test_loss, test_acc = ann.evaluate(test_images, test_labels)
print(f"Test accuracy: {test_acc}")

# Plot training history
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
plt.show()
```

Epoch 1/10
 1563/1563 [=====] - 9s 5ms/step - loss: 1.8968 - accuracy: 0.3092 - val_loss: 1.8106 - val_accuracy: 0.3435
 Epoch 2/10
 1563/1563 [=====] - 8s 5ms/step - loss: 1.7199 - accuracy: 0.3832 - val_loss: 1.6582 - val_accuracy: 0.4110
 Epoch 3/10
 1563/1563 [=====] - 8s 5ms/step - loss: 1.6422 - accuracy: 0.4098 - val_loss: 1.6291 - val_accuracy: 0.4232
 Epoch 4/10
 1563/1563 [=====] - 8s 5ms/step - loss: 1.5831 - accuracy: 0.4316 - val_loss: 1.5516 - val_accuracy: 0.4501
 Epoch 5/10
 1563/1563 [=====] - 8s 5ms/step - loss: 1.5467 - accuracy: 0.4464 - val_loss: 1.5464 - val_accuracy: 0.4532
 Epoch 6/10
 1563/1563 [=====] - 8s 5ms/step - loss: 1.5091 - accuracy: 0.4582 - val_loss: 1.5096 - val_accuracy: 0.4586
 Epoch 7/10
 1563/1563 [=====] - 8s 5ms/step - loss: 1.4828 - accuracy: 0.4683 - val_loss: 1.4898 - val_accuracy: 0.4718
 Epoch 8/10
 1563/1563 [=====] - 8s 5ms/step - loss: 1.4608 - accuracy: 0.4751 - val_loss: 1.4794 - val_accuracy: 0.4711
 Epoch 9/10
 1563/1563 [=====] - 8s 5ms/step - loss: 1.4388 - accuracy: 0.4843 - val_loss: 1.4736 - val_accuracy: 0.4776
 Epoch 10/10
 1563/1563 [=====] - 8s 5ms/step - loss: 1.4218 - accuracy: 0.4906 - val_loss: 1.4666 - val_accuracy: 0.4774
 313/313 [=====] - 0s 1ms/step - loss: 1.4666 - accuracy: 0.4774
 Test accuracy: 0.477400004863739



Results

In the context of this comparison, the deeper neural network model demonstrated a higher accuracy score when measured against the shallow network model. This outcome aligns with the typical trend observed in neural network architectures, emphasizing the advantages of leveraging deeper architectures for enhanced learning and better performance in learning complex patterns from the data.

Model 2: Convolutional Neural Networks (CNNs)

Having explored Artificial Neural Networks (ANNs), our attention now shifts to Convolutional Neural Networks (CNNs). These specialized neural networks are tailored for processing structured grid-like data, particularly prevalent in image and video recognition tasks. CNNs utilize convolutional layers, filters, and pooling layers to extract and learn hierarchical representations of features within images. By doing so, they effectively capture spatial hierarchies and patterns present in visual data.

CNNs excel in scenarios where understanding local patterns and spatial relationships is crucial. Their effectiveness in image recognition, object detection, and various computer vision tasks stems from their ability to automatically learn features from raw pixel data. Moreover, CNNs leverage parameter sharing across the input, significantly reducing the number of parameters. This parameter sharing enables efficient learning from high-dimensional inputs like images, contributing to both performance and computational efficiency.

Similar to the experimentation with ANNs, we will explore the impact of different optimizers within CNN architectures to discern their performance differences. Optimizers such as Adam, SGD, RMSProp, and others will be tested to observe their influence on the CNN's learning dynamics, convergence, and overall performance in image-related tasks.

Optimizers

```
In [8]: # Create CNN model
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),

    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax'), # 10 output classes for CIFAR-10
])

# Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

```

# Train the model
history = model.fit(train_images, train_labels, epochs=10,
                    validation_data=(test_images, test_labels))

# Evaluate the model
test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f"Test accuracy: {test_acc}")

# Plot training history
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
plt.show()


```

WARNING:tensorflow:From c:\Users\rodri\AppData\Local\Programs\Python\Python39\lib\site-packages\keras\src\layers\pooling\max_pooling2d.py:161: The name tf.nn.max_pool is deprecated. Please use tf.nn.max_pool2d instead.

```

Epoch 1/10
1563/1563 [=====] - 11s 7ms/step - loss: 1.4850 - accuracy:
0.4606 - val_loss: 1.2637 - val_accuracy: 0.5401
Epoch 2/10
1563/1563 [=====] - 10s 7ms/step - loss: 1.1126 - accuracy:
0.6079 - val_loss: 1.0450 - val_accuracy: 0.6294
Epoch 3/10
1563/1563 [=====] - 10s 7ms/step - loss: 0.9704 - accuracy:
0.6606 - val_loss: 0.9570 - val_accuracy: 0.6607
Epoch 4/10
1563/1563 [=====] - 10s 7ms/step - loss: 0.8827 - accuracy:
0.6910 - val_loss: 0.9103 - val_accuracy: 0.6811
Epoch 5/10
1563/1563 [=====] - 10s 7ms/step - loss: 0.8063 - accuracy:
0.7181 - val_loss: 0.8600 - val_accuracy: 0.6976
Epoch 6/10
1563/1563 [=====] - 10s 7ms/step - loss: 0.7505 - accuracy:
0.7366 - val_loss: 0.8690 - val_accuracy: 0.6974
Epoch 7/10
1563/1563 [=====] - 11s 7ms/step - loss: 0.7014 - accuracy:
0.7506 - val_loss: 0.8765 - val_accuracy: 0.6990
Epoch 8/10
1563/1563 [=====] - 10s 7ms/step - loss: 0.6549 - accuracy:
0.7726 - val_loss: 0.8426 - val_accuracy: 0.7154
Epoch 9/10
1563/1563 [=====] - 10s 7ms/step - loss: 0.6090 - accuracy:
0.7853 - val_loss: 0.8999 - val_accuracy: 0.7006
Epoch 10/10
1563/1563 [=====] - 10s 7ms/step - loss: 0.5727 - accuracy:
0.7970 - val_loss: 0.8672 - val_accuracy: 0.7134
313/313 [=====] - 1s 2ms/step - loss: 0.8672 - accuracy: 0.7
134
Test accuracy: 0.7134000062942505

```



```
In [9]: # Create CNN model
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),

    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax'), # 10 output classes for CIFAR-10
])


# Compile the model
model.compile(optimizer='sgd',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(train_images, train_labels, epochs=10,
                    validation_data=(test_images, test_labels))

# Evaluate the model
test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f"Test accuracy: {test_acc}")

# Plot training history
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
plt.show()
```

Epoch 1/10
 1563/1563 [=====] - 11s 7ms/step - loss: 2.0708 - accuracy: 0.2409 - val_loss: 1.8409 - val_accuracy: 0.3339
 Epoch 2/10
 1563/1563 [=====] - 10s 7ms/step - loss: 1.6765 - accuracy: 0.3966 - val_loss: 1.5816 - val_accuracy: 0.4287
 Epoch 3/10
 1563/1563 [=====] - 10s 6ms/step - loss: 1.4992 - accuracy: 0.4625 - val_loss: 1.5324 - val_accuracy: 0.4678
 Epoch 4/10
 1563/1563 [=====] - 10s 6ms/step - loss: 1.3859 - accuracy: 0.5073 - val_loss: 1.3369 - val_accuracy: 0.5180
 Epoch 5/10
 1563/1563 [=====] - 10s 7ms/step - loss: 1.2930 - accuracy: 0.5433 - val_loss: 1.2802 - val_accuracy: 0.5444
 Epoch 6/10
 1563/1563 [=====] - 10s 6ms/step - loss: 1.2165 - accuracy: 0.5700 - val_loss: 1.2965 - val_accuracy: 0.5517
 Epoch 7/10
 1563/1563 [=====] - 10s 6ms/step - loss: 1.1503 - accuracy: 0.5967 - val_loss: 1.1826 - val_accuracy: 0.5856
 Epoch 8/10
 1563/1563 [=====] - 10s 6ms/step - loss: 1.0928 - accuracy: 0.6181 - val_loss: 1.1001 - val_accuracy: 0.6100
 Epoch 9/10
 1563/1563 [=====] - 10s 6ms/step - loss: 1.0438 - accuracy: 0.6347 - val_loss: 1.1159 - val_accuracy: 0.5987
 Epoch 10/10
 1563/1563 [=====] - 10s 6ms/step - loss: 0.9988 - accuracy: 0.6533 - val_loss: 1.0537 - val_accuracy: 0.6319
 313/313 [=====] - 1s 2ms/step - loss: 1.0537 - accuracy: 0.6319
 Test accuracy: 0.6319000124931335



```
In [10]: # Create CNN model
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),

    layers.Flatten(),
    layers.Dense(64, activation='relu'),
```

```

layers.Dense(10, activation='softmax'), # 10 output classes for CIFAR-10
])

# Compile the model
model.compile(optimizer='rmsprop',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(train_images, train_labels, epochs=10,
                    validation_data=(test_images, test_labels))

# Evaluate the model
test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f"Test accuracy: {test_acc}")


# Plot training history
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
plt.show()

```

```

Epoch 1/10
1563/1563 [=====] - 11s 7ms/step - loss: 1.5557 - accuracy:
0.4384 - val_loss: 1.3610 - val_accuracy: 0.4962
Epoch 2/10
1563/1563 [=====] - 10s 7ms/step - loss: 1.1467 - accuracy:
0.5968 - val_loss: 1.1247 - val_accuracy: 0.5964
Epoch 3/10
1563/1563 [=====] - 10s 7ms/step - loss: 0.9815 - accuracy:
0.6563 - val_loss: 0.9868 - val_accuracy: 0.6583
Epoch 4/10
1563/1563 [=====] - 10s 7ms/step - loss: 0.8702 - accuracy:
0.7005 - val_loss: 1.0705 - val_accuracy: 0.6333
Epoch 5/10
1563/1563 [=====] - 11s 7ms/step - loss: 0.7896 - accuracy:
0.7270 - val_loss: 0.9581 - val_accuracy: 0.6720
Epoch 6/10
1563/1563 [=====] - 10s 7ms/step - loss: 0.7225 - accuracy:
0.7508 - val_loss: 0.9885 - val_accuracy: 0.6830
Epoch 7/10
1563/1563 [=====] - 10s 7ms/step - loss: 0.6723 - accuracy:
0.7664 - val_loss: 0.9011 - val_accuracy: 0.7078
Epoch 8/10
1563/1563 [=====] - 10s 7ms/step - loss: 0.6267 - accuracy:
0.7839 - val_loss: 0.9414 - val_accuracy: 0.7004
Epoch 9/10
1563/1563 [=====] - 10s 7ms/step - loss: 0.5884 - accuracy:
0.7976 - val_loss: 1.1393 - val_accuracy: 0.6830
Epoch 10/10
1563/1563 [=====] - 10s 7ms/step - loss: 0.5589 - accuracy:
0.8087 - val_loss: 1.0724 - val_accuracy: 0.6845
313/313 [=====] - 1s 2ms/step - loss: 1.0724 - accuracy: 0.6
845
Test accuracy: 0.684499979019165

```



Results

Upon evaluating the performance of optimizers in a CNN model, the Adam optimizer emerges as the most effective among the three—Adam, SGD, and RMSProp—utilized in the image classification task, achieving an accuracy score of 71.34%.

Layers

```
In [11]: # Create a Shallow Networks model
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),

    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax'), # 10 output classes for CIFAR-10
])


# Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(train_images, train_labels, epochs=10,
                    validation_data=(test_images, test_labels))

# Evaluate the model
test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f"Test accuracy: {test_acc}")

# Plot training history
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
plt.show()
```


Epoch 1/10
 1563/1563 [=====] - 14s 9ms/step - loss: 1.3470 - accuracy: 0.5200 - val_loss: 1.0885 - val_accuracy: 0.6154
 Epoch 2/10
 1563/1563 [=====] - 13s 9ms/step - loss: 0.9663 - accuracy: 0.6641 - val_loss: 0.9634 - val_accuracy: 0.6628
 Epoch 3/10
 1563/1563 [=====] - 13s 9ms/step - loss: 0.8200 - accuracy: 0.7145 - val_loss: 0.9028 - val_accuracy: 0.6852
 Epoch 4/10
 1563/1563 [=====] - 13s 9ms/step - loss: 0.7072 - accuracy: 0.7552 - val_loss: 0.8980 - val_accuracy: 0.6964
 Epoch 5/10
 1563/1563 [=====] - 13s 9ms/step - loss: 0.6020 - accuracy: 0.7904 - val_loss: 0.9076 - val_accuracy: 0.7026
 Epoch 6/10
 1563/1563 [=====] - 13s 8ms/step - loss: 0.5072 - accuracy: 0.8230 - val_loss: 0.9314 - val_accuracy: 0.6998
 Epoch 7/10
 1563/1563 [=====] - 13s 8ms/step - loss: 0.4254 - accuracy: 0.8522 - val_loss: 1.0397 - val_accuracy: 0.6972
 Epoch 8/10
 1563/1563 [=====] - 13s 9ms/step - loss: 0.3474 - accuracy: 0.8774 - val_loss: 1.1506 - val_accuracy: 0.6861
 Epoch 9/10
 1563/1563 [=====] - 13s 9ms/step - loss: 0.2802 - accuracy: 0.9022 - val_loss: 1.2322 - val_accuracy: 0.6908
 Epoch 10/10
 1563/1563 [=====] - 13s 8ms/step - loss: 0.2268 - accuracy: 0.9213 - val_loss: 1.3502 - val_accuracy: 0.6883
 313/313 [=====] - 1s 2ms/step - loss: 1.3502 - accuracy: 0.6883
 Test accuracy: 0.6883000135421753



```
In [22]: # Create a Deep Networks model
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.Conv2D(128, (3, 3), activation='relu'),

    layers.Flatten(),
```

```

layers.Dense(64, activation='relu'),
layers.Dense(10, activation='softmax'), # 10 output classes for CIFAR-10
])

# Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Train the model
history = model.fit(train_images, train_labels, epochs=10,
                    validation_data=(test_images, test_labels))

# Evaluate the model
test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f"Test accuracy: {test_acc}")


# Plot training history
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
plt.show()

```

```

Epoch 1/10
1563/1563 [=====] - 13s 8ms/step - loss: 1.5578 - accuracy:
0.4273 - val_loss: 1.3352 - val_accuracy: 0.5039
Epoch 2/10
1563/1563 [=====] - 12s 8ms/step - loss: 1.1828 - accuracy:
0.5768 - val_loss: 1.1015 - val_accuracy: 0.6020
Epoch 3/10
1563/1563 [=====] - 12s 8ms/step - loss: 1.0156 - accuracy:
0.6420 - val_loss: 1.0412 - val_accuracy: 0.6345
Epoch 4/10
1563/1563 [=====] - 12s 8ms/step - loss: 0.9093 - accuracy:
0.6813 - val_loss: 0.9930 - val_accuracy: 0.6502
Epoch 5/10
1563/1563 [=====] - 12s 8ms/step - loss: 0.8264 - accuracy:
0.7078 - val_loss: 0.9513 - val_accuracy: 0.6719
Epoch 6/10
1563/1563 [=====] - 12s 8ms/step - loss: 0.7620 - accuracy:
0.7338 - val_loss: 0.8997 - val_accuracy: 0.6880
Epoch 7/10
1563/1563 [=====] - 12s 8ms/step - loss: 0.7004 - accuracy:
0.7533 - val_loss: 0.8904 - val_accuracy: 0.6966
Epoch 8/10
1563/1563 [=====] - 12s 8ms/step - loss: 0.6504 - accuracy:
0.7711 - val_loss: 0.8894 - val_accuracy: 0.6967
Epoch 9/10
1563/1563 [=====] - 12s 7ms/step - loss: 0.6077 - accuracy:
0.7836 - val_loss: 0.9343 - val_accuracy: 0.6923
Epoch 10/10
1563/1563 [=====] - 12s 8ms/step - loss: 0.5636 - accuracy:
0.7996 - val_loss: 0.8934 - val_accuracy: 0.7138
313/313 [=====] - 1s 3ms/step - loss: 0.8934 - accuracy: 0.7
138
Test accuracy: 0.7138000130653381

```



Results

In this comparison, the deeper neural network model exhibited superior accuracy compared to the shallow network model. This aligns with the commonly observed trend in neural network architectures, highlighting the benefits of utilizing deeper structures. Deeper architectures excel in learning intricate patterns from data, resulting in enhanced learning capabilities and improved performance.